



Jeff Knupp

PYTHON PROGRAMMER

[BLOG](#) [ABOUT](#) [ARCHIVES](#) [TUTORING](#) [BOOK](#)

Everything I know about Python...

Learn to Write Pythonic Code!

Check out the book *Writing Idiomatic Python!*

Looking for Python Tutoring? Remote and local (NYC) slots still available! Email me at jeff@jeffknupp.com for more info.

Drastically Improve Your Python: Understanding Python's Execution Model

Those new to Python are often surprised by the behavior of their own code. They expect *A* but, seemingly for no reason, *B* happens instead. The root cause of many of these "surprises" is confusion about the Python execution model. It's the sort of thing that, if it's explained to you once, a number of Python concepts that seemed hazy before become crystal clear. It's also really difficult to just "figure out" on your own, as it requires a fundamental shift in thinking about core language concepts like variables, objects, and functions.

In this post, I'll help you understand what's happening behind the scenes when you do common things like creating a variable or calling a function. As a result, you'll write cleaner, more comprehensible code. You'll also become a better (and faster) code reader. All that's necessary is to forget everything you know about programming...

"Everything is an object?"

When most people first hear that in Python, "everything is an object", it triggers flashbacks to languages like Java where everything the *user* writes is encapsulated in an object. Others assume this means that in the implementation of the Python interpreter, everything is implemented as objects. The first interpretation is wrong; the second is true but not particularly interesting (for our purposes). What the phrase actually refers to is the fact that all "things", be they values, classes, functions, object instances (obviously), and almost every other language construct is conceptually an object.

What does it mean for everything to be an object? It means all of the "things" mentioned above have all the properties we usually associate with objects (in the object oriented sense); types have member functions, functions have attributes, modules can be passed as arguments, etc. And it has important implications with regards to how assignment in Python works.

A feature of the Python interpreter that often confuses beginners is what happens when `print()` is called on a "variable" assigned to a user-defined object (I'll explain the quotes in a second). With built-in types, a proper value is usually printed, like when calling `print()` on `strings` and `ints`. For simple, user-defined classes, though, the interpreter spits out some odd looking string like:

```
>>> class Foo(): pass
>>> foo = Foo()
```

```
>>> print(foo)
<__main__.Foo object at 0xd3adb33f>
```

`print()` is supposed to print the value of a "variable", right? So why is it printing that garbage?

To answer that, we need to understand what `foo` actually represents in Python. Most other languages would call it a variable. Indeed, many Python articles would refer to `foo` as a variable, but really only as a shorthand notation.

In languages like C, `foo` represents storage for "stuff". If we wrote

```
int foo = 42;
```

it would be correct to say that the integer variable `foo` contained the value `42`. That is, *variables are a sort of container for values*.

And now for something completely different...

In Python, this isn't the case. When we say:

```
>>> foo = Foo()
```

it would be wrong to say that `foo` "contained" a `Foo` object. Rather, `foo` is a *name* with a *binding* to the *object* created by `Foo()`. The portion of the right hand side of the equals sign creates an object. Assigning `foo` to that object merely says "I want to be able to refer to this object as `foo`". **Instead of variables (in the classic sense), Python has names and bindings.**

So when we printed `foo` earlier, what the interpreter was showing us was the address in memory where the object that `foo` is bound to is stored. This isn't as useless as it sounds. If you're in the interpreter and want to see if two names are bound to the same object, you can do a quick-and-dirty check by printing them and comparing the addresses. If they match, they're

bound to the same object; if not, their bound to different objects. Of course, the idiomatic way to check if two names are bound to the same object is to use `is`

If we continued our example and wrote

```
>>> baz = foo
```

we should read this as "Bind the name `baz` to the same object `foo` is bound to (whatever that may be)." It should be clear, then why the following happens

```
>>> baz.some_attribute
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
AttributeError: 'Foo' object has no attribute 'some_attribute'
>>> foo.some_attribute = 'set from foo'
>>> baz.some_attribute
'set from foo'
```

Changing the object in some way using `foo` will also be reflected in `baz`: they are both bound to the same underlying object.

What's in a name...

`names` in Python are not unlike names in the real world. If my wife calls me "Jeff", my dad calls me "Jeffrey", and my boss calls me "Idiot", it doesn't fundamentally change *me*. If my boss decides to call me "Captain Programming," great, but it still hasn't changed anything about me. It does mean, however, that if my wife kills "Jeff" (and who could blame her), "Captain Programming" is also dead. Likewise, in Python binding a name to an object doesn't change it. Changing some property of the object, however, will be reflected in all other names bound to that object.

Everything really *is* an object. I swear.

Here, a question arises: How do we know that the thing on the right hand side of the equals sign will always be an object we can bind a name to? What about

```
>>> foo = 10
```

or

```
>>> foo = "Hello World!"
```

Now is when "everything is an object" pays off. Anything you can (legally) place on the right hand side of the equals sign is (or creates) an object in Python. Both `10` and `Hello World` are objects. Don't believe me? Check for yourself

```
>>> foo = 10
>>> print(foo.__add__)
<method-wrapper '__add__' of int object at 0x8502c0>
```

If `10` was actually just the number '10', it probably wouldn't have an `__add__` attribute (or any attributes at all).

In fact, we can see all the attributes `10` has using the `dir()` function:

```
>>> dir(10)
['_abs_', '__add__', '__and__', '__class__', '__cmp__', '__coerce__', '__delattr__',
 '__div__', '__divmod__', '__doc__', '__float__', '__floordiv__', '__format__',
 '__getattr__', '__getnewargs__', '__hash__', '__hex__', '__index__',
 '__init__', '__int__', '__invert__', '__long__', '__lshift__', '__mod__',
 '__mul__', '__neg__', '__new__', '__nonzero__', '__oct__', '__or__',
 '__pos__', '__pow__', '__radd__', '__rand__', '__rdiv__', '__rdivmod__',
 '__reduce__', '__reduce_ex__', '__repr__', '__rfloordiv__', '__rlshift__',
 '__rmod__', '__rmul__', '__ror__', '__rpow__', '__rrshift__', '__rshift__',
 '__rsub__', '__rtruediv__', '__rxor__', '__setattr__', '__sizeof__', '__str__',
 '__sub__', '__subclasshook__', '__truediv__', '__trunc__', '__xor__',
 'bit_length', 'conjugate', 'denominator', 'imag', 'numerator', 'real']
```

With all those attributes and member functions, I think it's safe to say `10` is an object.

Since everything in Python is essentially names bound to objects, we can do silly (but interesting) stuff like this:

```
>>> import datetime
>>> import imp
>>> datetime.datetime.now()
datetime.datetime(2013, 02, 14, 02, 53, 59, 608842)
>>> class PartyTime():
...     def __call__(self, *args):
...         imp.reload(datetime)
...         value = datetime.datetime(*args)
...         datetime.datetime = self
...         return value
...
...     def __getattr__(self, value):
...         if value == 'now':
...             return lambda: print('Party Time!')
...         else:
...             imp.reload(datetime)
...             value = getattr(datetime.datetime, value)
...             datetime.datetime = self
...             return value
>>> datetime.datetime = PartyTime()
>>> datetime.datetime.now()
Party Time!
>>> today = datetime.datetime(2013, 2, 14)
>>> print(today)
2013-02-14 00:00:00
>>> print(today.timestamp())
1360818000.0
```

`datetime.datetime` is just a name (that happens to be bound to an object representing the `datetime` class). We can rebind it to whatever we please. In the example above, we bind the `datetime` attribute of the `datetime` module to our new class, `PartyTime`. Any call to the `datetime.datetime` constructor returns a valid `datetime` object. In fact, the class is

indistinguishable from the real `datetime.datetime` class. Except, that is, for the fact that if you call `datetime.datetime.now()` it always prints out 'Party Time!'.

Obviously this is a silly example, but hopefully it gives you some insight into what is possible when you fully understand and make use of Python's execution model. At this point, though, we've only changed bindings associated with a name. What about changing the object itself?

Two types of objects

It turns out Python has two flavors of objects: `mutable` and `immutable`. The value of mutable objects can be changed after they are created. The value of immutable objects cannot be. A `list` is a mutable object. You can create a list, append some values, and the list is updated in place. A `string` is immutable. Once you create a string, you can't change its value.

I know what you're thinking: "Of course you can change the value of a string, I do it all the time in my code!" When you "change" a string, you're actually rebinding it to a newly created string object. The original object remains unchanged, even though it's possible that nothing refers to it anymore.

See for yourself:

```
>>> a = 'foo'
>>> a
'foo'
>>> b = a
>>> a += 'bar'
>>> a
'foobar'
>>> b
'foo'
```

Even though we're using `+=` and it *seems* that we're modifying the string, we really just get a new one containing the result of the change. This is why you may hear people say, "string concatenation is slow.". It's because concatenating strings must allocate memory for a new string and copy the contents, while appending to a list (in most cases) requires no allocation.

Immutable objects are fundamentally expensive to "change", because doing so involves creating a copy. Changing mutable objects is cheap.

Immutable object weirdness

When I said the value of immutable objects can't change after they're created, it wasn't the whole truth. A number of containers in Python, such as `tuple`, are immutable. The value of a `tuple` can't be changed after it is created. But the "value" of a tuple is conceptually just a sequence of names with unchangeable bindings to objects. The key thing to note is that the *bindings* are unchangeable, not the objects they are bound to.

This means the following is perfectly legal:

```
>>> class Foo():
...     def __init__(self):
...         self.value = 0
...     def __str__(self):
...         return str(self.value)
...     def __repr__(self):
...         return str(self.value)
...
>>> f = Foo()
>>> print(f)
0
>>> foo_tuple = (f, f)
>>> print(foo_tuple)
(0, 0)
>>> foo_tuple[0] = 100
Traceback (most recent call last):
File "<stdin>", line 1, in <module>
TypeError: 'tuple' object does not support item assignment
>>> f.value = 999
>>> print(f)
999
>>> print(foo_tuple)
(999, 999)
```


When we try to change an element of the tuple directly, we get a `TypeError` telling us that (once created), `tuples` can't be assigned to. But changing the underlying object has the effect of "changing" the value of the `tuple`. This is a subtle point, but nonetheless important: the "value" of an immutable object *can't* change, but its constituent objects *can*.

Function calls

If variables are just `names` bound to objects, what happens when we pass them as arguments to a function? The truth is, we aren't really passing all that much. Take a look at this code:

```
def add_to_tree(root, value_string):
    """Given a string of characters `value_string`, create or update a
    series of dictionaries where the value at each level is a dictionary of
    the characters that have been seen following the current character.

    Example:
    >>> my_string = 'abc'
    >>> tree = {}
    >>> add_to_tree(tree, my_string)
    >>> print(tree['a']['b'])
    {'c': {}}
    >>> add_to_tree(tree, 'abd')
    >>> print(tree['a']['b'])
    {'c': {}, 'd': {}}
    >>> print(tree['a']['d'])
    KeyError 'd'
    """

    for character in value_string:
        root = root.setdefault(character, {})
```

We're essentially creating an auto-vivifying dictionary that operates like a trie. Notice that we change the `root` parameter in the `for` loop. And yet after the function call completes, `tree` is still the same dictionary with some updates. It is *not* the last

value of `root` in the function call. So in one sense `tree` is being updated; in another sense it's not.

To make sense of this, consider what the `root` parameter actually is: a *new* binding to the object referred to by the name passed in as the `root` parameter. In the case of our example, `root` is a name initially bound to the same object as `tree`. It is *not* `tree` itself, which explains why changing `root` to a new dictionary in the function leaves `tree` unchanged. As you'll recall, assigning `root` to `root.setdefault(character, {})` merely rebinds `root` to the object created by the `root.setdefault(character, {})` statement.

Here's another, more straightforward, example:

```
def list_changer(input_list):
    input_list[0] = 10

    input_list = range(1, 10)
    print(input_list)
    input_list[0] = 10
    print(input_list)

>>> test_list = [5, 5, 5]
>>> list_changer(test_list)
[1, 2, 3, 4, 5, 6, 7, 8, 9]
[10, 2, 3, 4, 5, 6, 7, 8, 9]
>>> print test_list
[10, 5, 5]
```

Our first statement *does* change the value of the underlying list (as we can see in the last line printed). However, once we rebind the name `input_list` by saying `input_list = range(1, 10)`, **we're now referring to a completely different object**. We basically said "bind the name `input_list` to this new list." After that line, we have no way of referring to the original `input_list` parameter again.

By now, you should have a clear understanding of how binding a name works. There's just one more item to take care of.

Blocks and Scope

The concepts of `names`, `bindings`, and `objects` should be quite familiar at this point. What we haven't covered, though, is how the interpreter "finds" a name. To see what I mean, consider the following:

```
GLOBAL_CONSTANT = 42

def print_some_weird_calculation(value):
    number_of_digits = len(str(value))

    def print_formatted_calculation(result):
        print('{value} * {constant} = {result}'.format(value=value,
            constant=GLOBAL_CONSTANT, result=result))
        print('{} {}'.format('^' * number_of_digits, '++'))
        print('\nKey: ^ points to your number, + points to constant')

    print_formatted_calculation(value * GLOBAL_CONSTANT)

>>> print_some_weird_calculation(123)
123 * 42 = 5166
^^^  ++

Key: ^ points to your number, + points to constant
```

This is a contrived example, but a couple of things should jump out at you. First, how does the `print_formatted_calculation` function have access to `value` and `number_of_digits` even though they were never passed as arguments? Second, how do both functions seem to have access to `GLOBAL_CONSTANT`?

The answer is all about `scope`. In Python, when a name is bound to an object, that name is only usable within the name's `scope`. The `scope` of a name is determined by the `block` in which it was created. A `block` is just a "block" of Python code that is executed as a single unit. The three most common types of blocks are modules, class definitions, and the bodies of functions. So the `scope` of a name is the innermost `block` in which it's defined.

Let's now return to the original question: how does the interpreter "find" what a name is bound to (or if it's even a valid name at all)? It begins by checking the `scope` of the innermost `block`. Then it checks the `scope` that contained the innermost `block`, then the `scope` that contained that, and so on.

In the `print_formatted_calculation` function, we reference `value`. This is resolved by first checking the `scope` of the innermost `block`, which in this case is the body of the function itself. When it doesn't find `value` defined there, it checks the `scope` that `print_formatted_calculation` was defined in. In our case, that's the body of the `print_some_weird_calculation` function. Here it does find the name `value`, and so it uses that binding and stops looking. The same is true for `GLOBAL_CONSTANT`, it just needs to look an extra level higher: the module (or script) level. Anything defined at this level is considered a `global` name. These are accessible from anywhere.

A few quick things to note. A name's `scope` extends to any blocks contained in the block where the name was defined, *unless the name is rebound in one of those blocks*. If `print_formatted_calculation` had the line `value = 3`, then the `scope` of the name `value` in `print_some_weird_calculation` would only be the body of that function. It's `scope` would not include `print_formatted_calculation`, since that `block` rebound the name.

Use this power wisely...

There are two keywords that can be used to tell the interpreter to **reuse a preexisting binding**. Every other time we bind a name, it binds that name to a new object, *but only in the current scope*. In the example above, if we rebound `value` in `print_formatted_calculation`, it would have no affect on the `value` in `print_some_weird_calculation`, which is `print_formatted_calculation`'s enclosing scope. With the following two keywords, we can actually affect the bindings outside our local scope.

`global my_variable` tells the interpreter to use the binding of the name `my_variable` in the top-most (or "global" scope). Putting `global my_variable` in a code block is a way of saying, "copy the binding of this global variable, or if you don't find it, create the name `my_variable` in the global scope." Similarly, the `nonlocal my_variable` statement instructs the interpreter to use the binding of the name `my_variable` defined in the nearest *enclosing* scope. This is a way to rebound a

name not defined in either the local or global scope. Without `nonlocal`, we would only be able to alter bindings in the local scope or the global scope. Unlike `global my_variable` however, if we use `nonlocal my_variable` then `my_variable` must already exist; it won't be created if it's not found.

To see this in action, let's write a quick example:

```
GLOBAL_CONSTANT = 42
print(GLOBAL_CONSTANT)
def outer_scope_function():
    some_value = hex(0x0)
    print(some_value)

    def inner_scope_function():
        nonlocal some_value
        some_value = hex(0xDEADBEEF)

    inner_scope_function()
    print(some_value)
    global GLOBAL_CONSTANT
    GLOBAL_CONSTANT = 31337

outer_scope_function()
print(GLOBAL_CONSTANT)

# Output:
# 42
# 0x0
# 0xdeadbeef
# 31337
```

By making use of `global` and `nonlocal`, we're able to use and change the existing binding of a name rather than merely assigning the name a new binding and losing the old one.

Summary

If you've made it to the end of this post, congratulations! Hopefully Python's execution model is much more clear. In a (much shorter) follow-up post, I'll show some examples of how we can make use of the fact that everything is an object in interesting ways. Until next time...

If you found this post useful, check out [Writing Idiomatic Python](#). It's filled with common Python idioms and code samples showing the right and wrong way to use them.

Posted on Feb 14, 2013 by Jeff Knupp



25 people like this. [Sign Up](#) to see what your friends like.

[« Write Cleaner Python: Use Exceptions](#)

Like this article?

Why not sign up for **Python Tutoring**? Sessions can be held remotely using Google+/Skype or in-person if you're in the NYC area. [Email jeff@jeffknupp.com](mailto:jeff@jeffknupp.com) if interested.

Sign up for the free jeffknupp.com email newsletter. Sent roughly once a month, it focuses on Python programming, scalable web development, and growing your freelance consultancy. And of course, you'll never be spammed, your privacy is protected, and you can opt out at any time.

* indicates required

Email Address *

Subscribe

